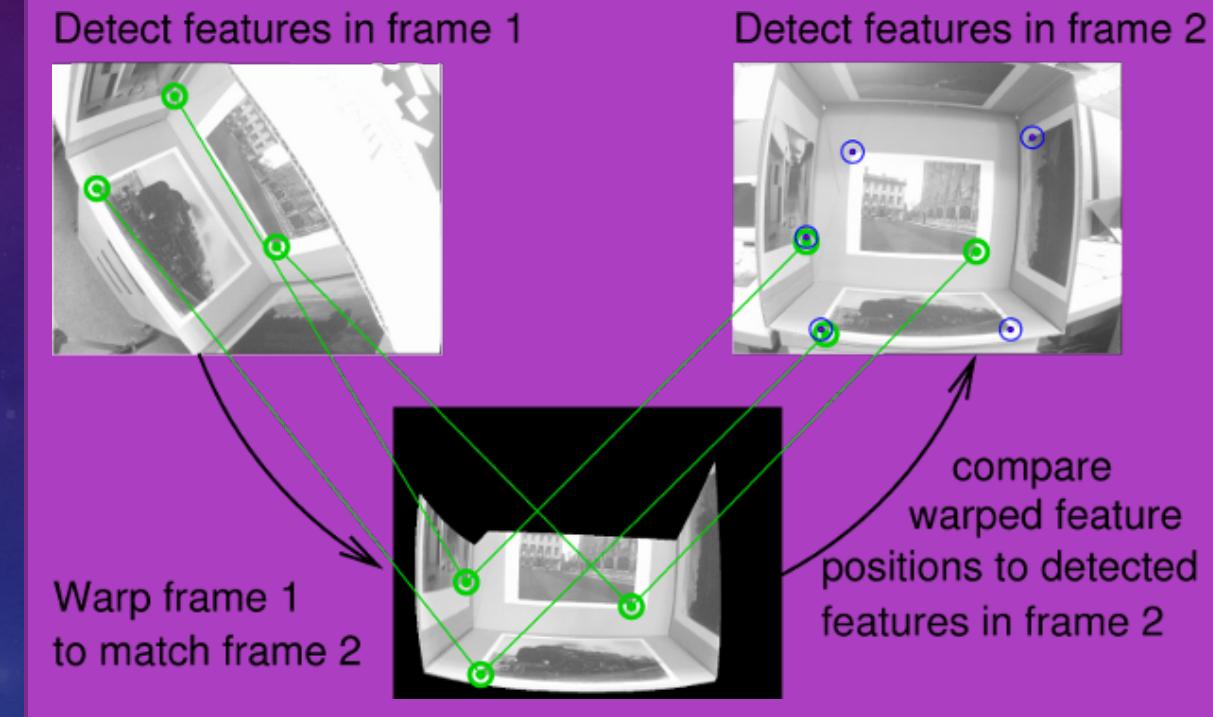


# ACCELERATED HARRIS CORNER DETECTOR

ANDREW TU  
EECE5640  
SPRING 2019

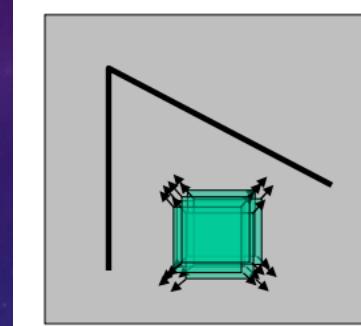
# CORNER DETECTORS

- Corners make good features to track
  - Rotation Invariant
  - Translation Invariant
- Useful for matching point between two images

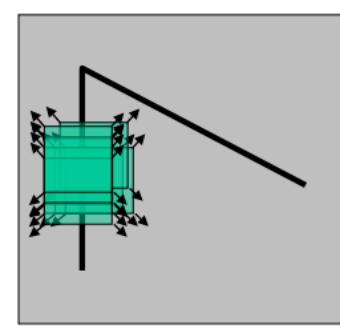


# HARRIS CORNER DETECTOR

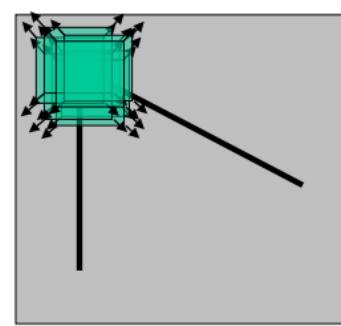
Find large changes in intensity in BOTH x and y direction



“flat” region:  
no change in  
all directions



“edge”:  
no change along  
the edge direction



“corner”:  
significant change  
in all directions

# HARRIS CORNER DETECTOR

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

$$R = \det M - k(\text{trace } M)^2$$

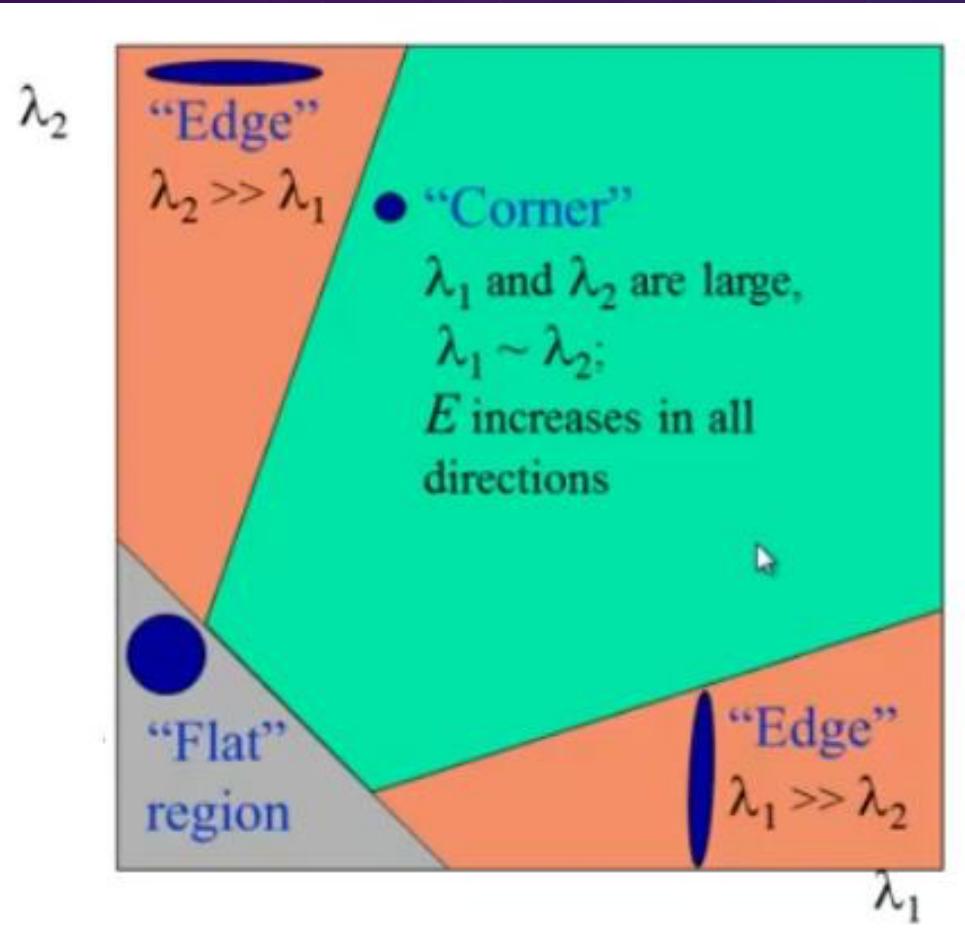
$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

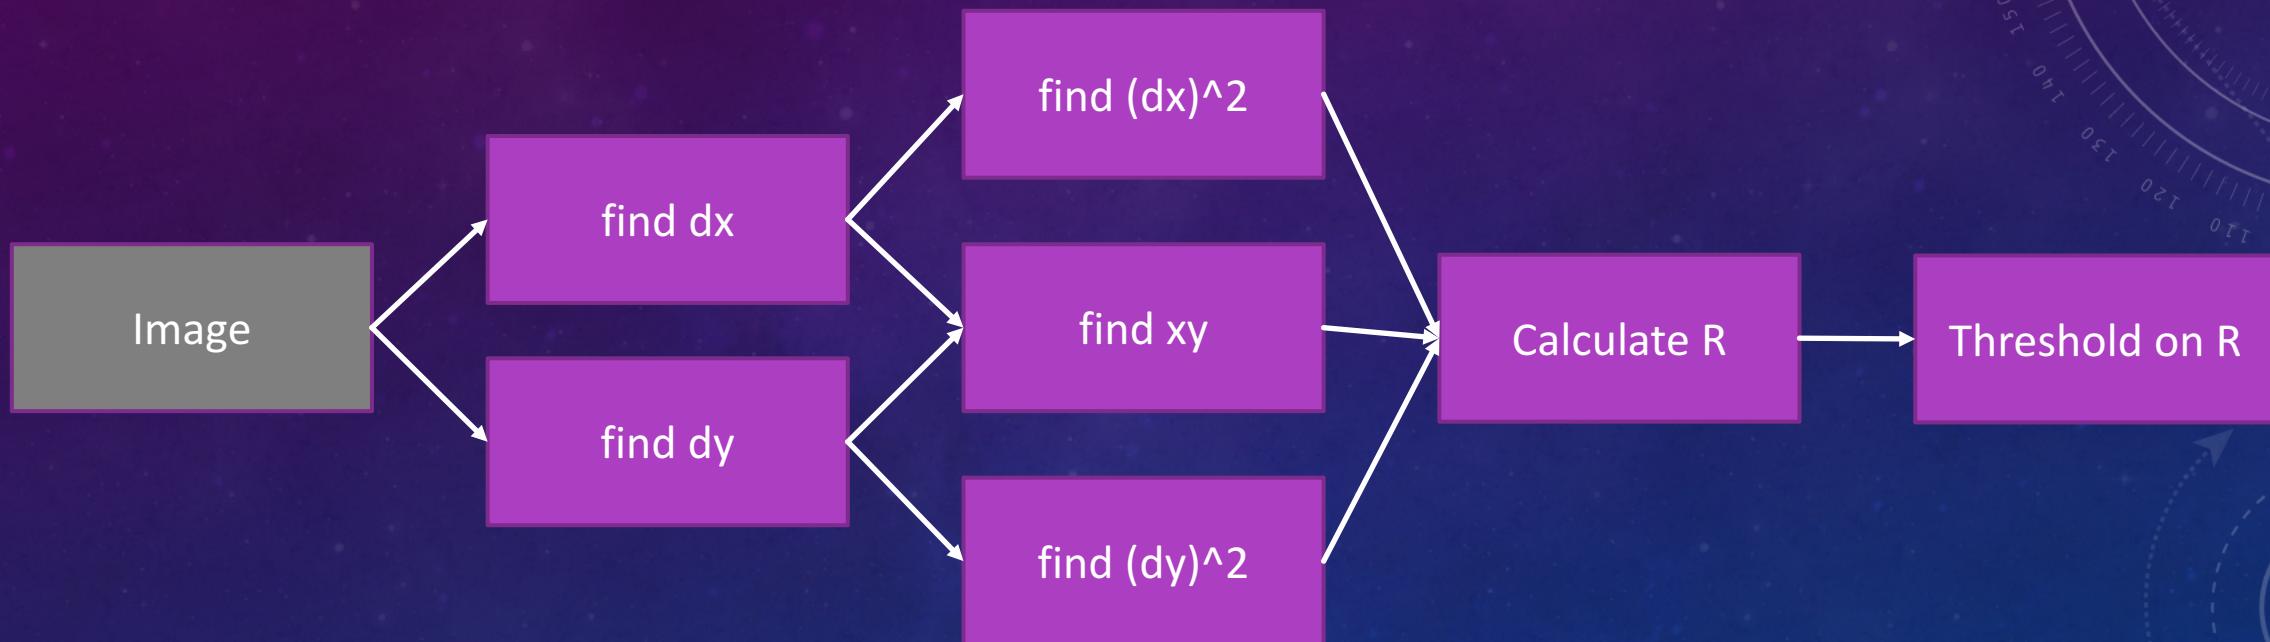
$$\lambda_1 = |I_x|^2$$

$$\lambda_2 = |I_y|^2$$

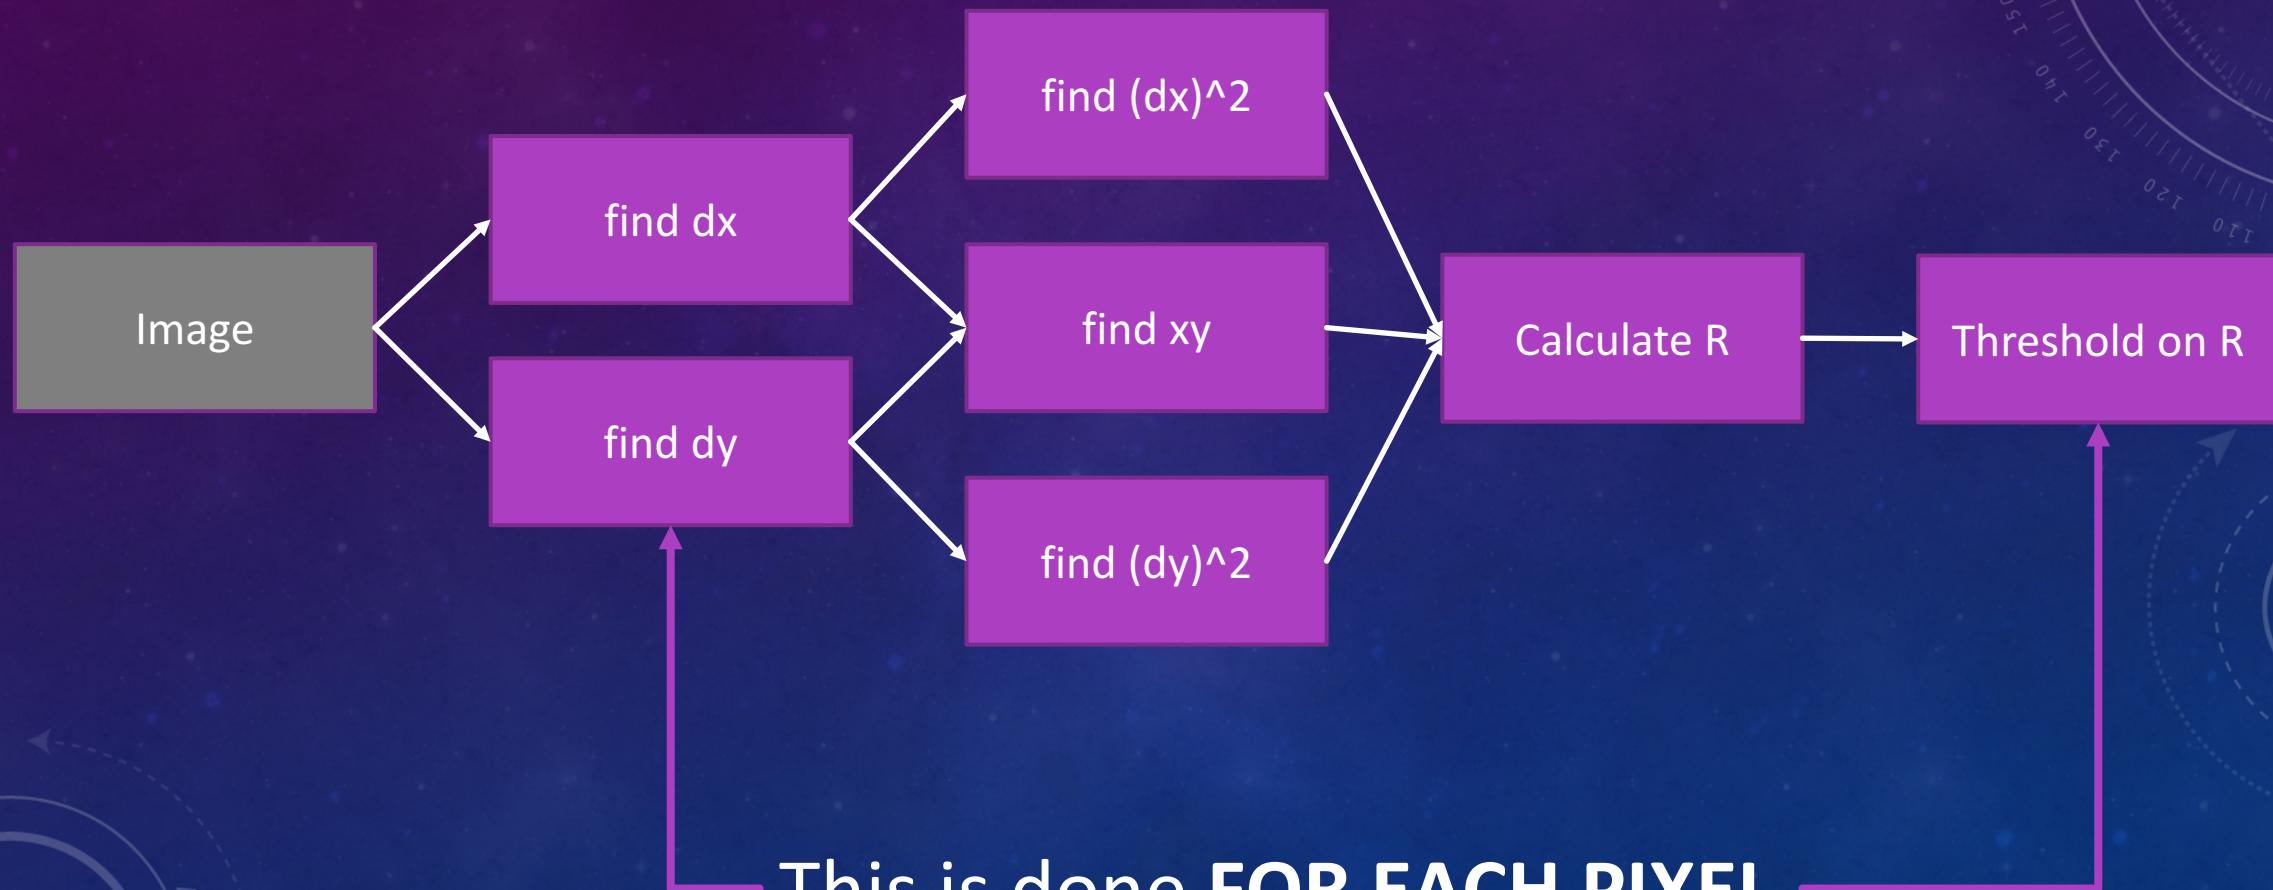
# HARRIS CORNER DETECTOR



# HARRIS CORNER DETECTOR PIPELINE

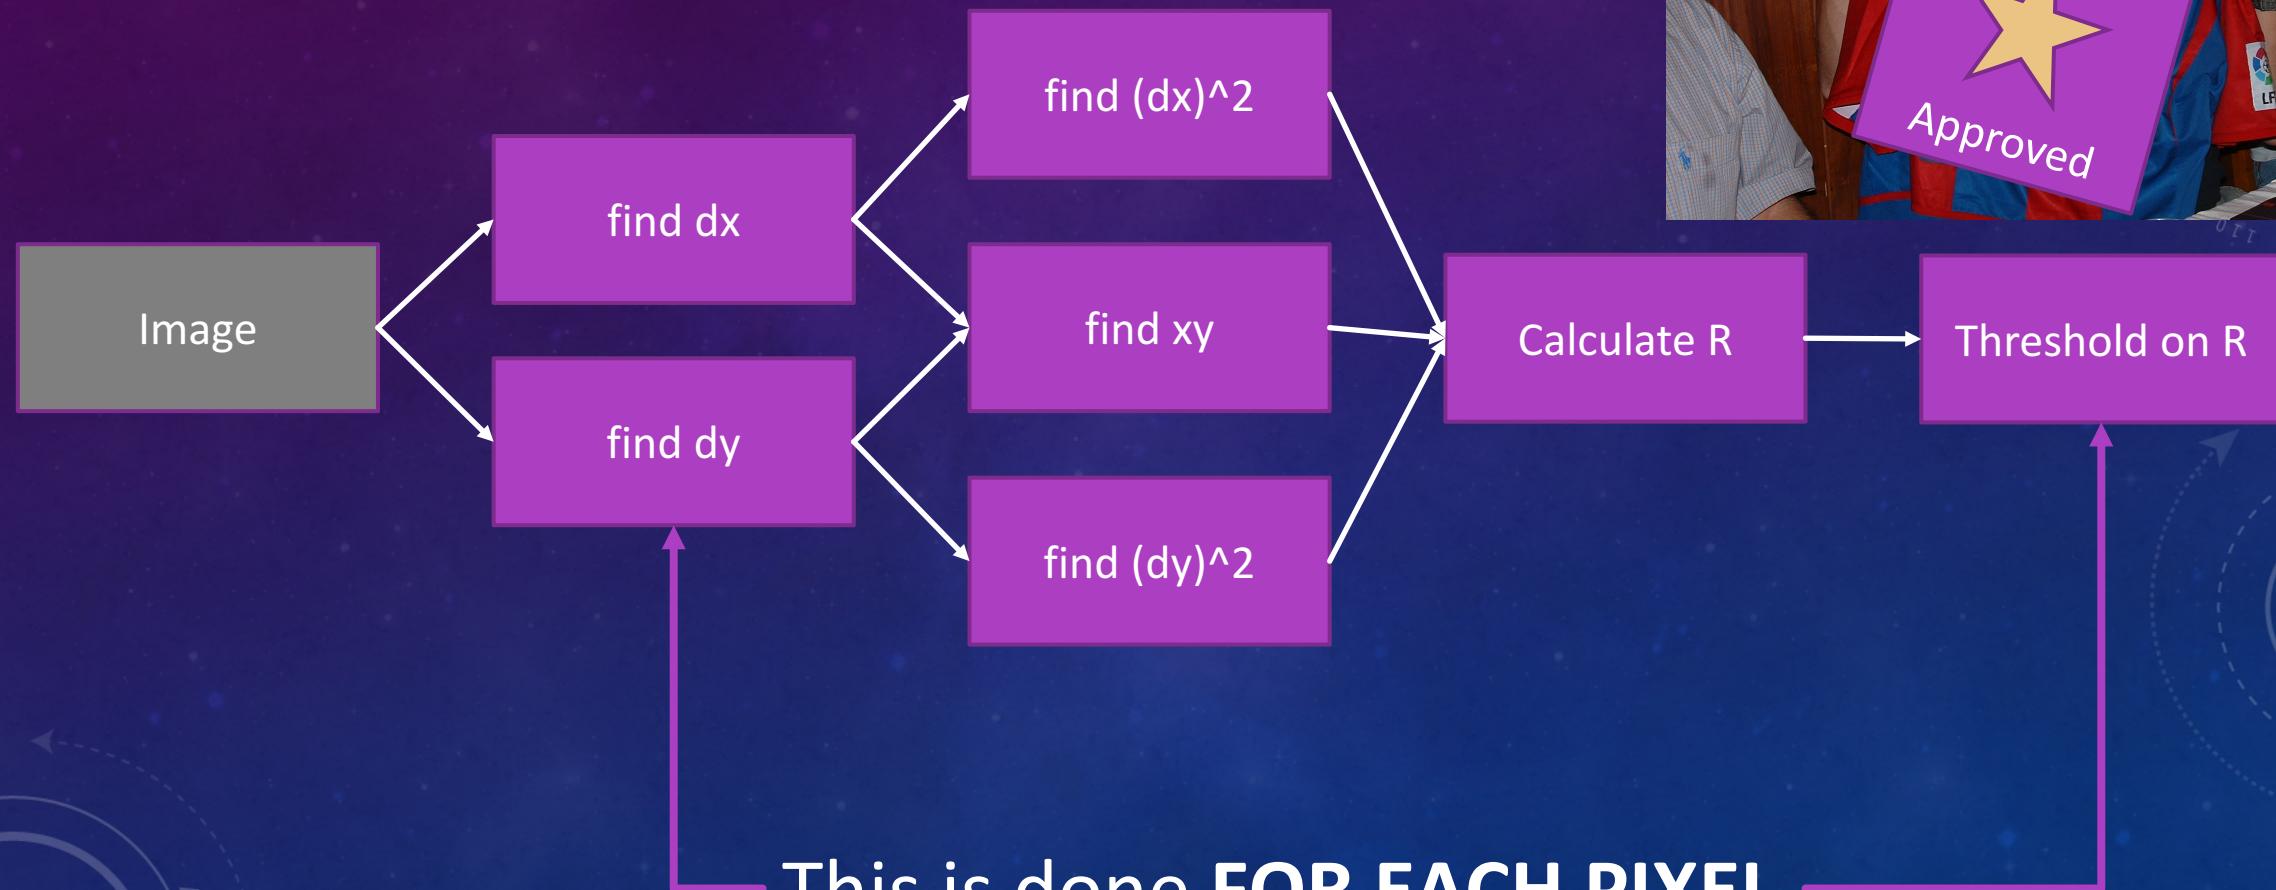


# HARRIS CORNER DETECTOR PIPELINE



This is done **FOR EACH PIXEL**

# HARRIS CORNER DETECTOR PIPELINE



# PROJECT GOALS

## Implementation Goals

1. Implement Harris Corner in C++
2. Accelerate with OpenMP
3. Accelerate with CUDA

## Analysis

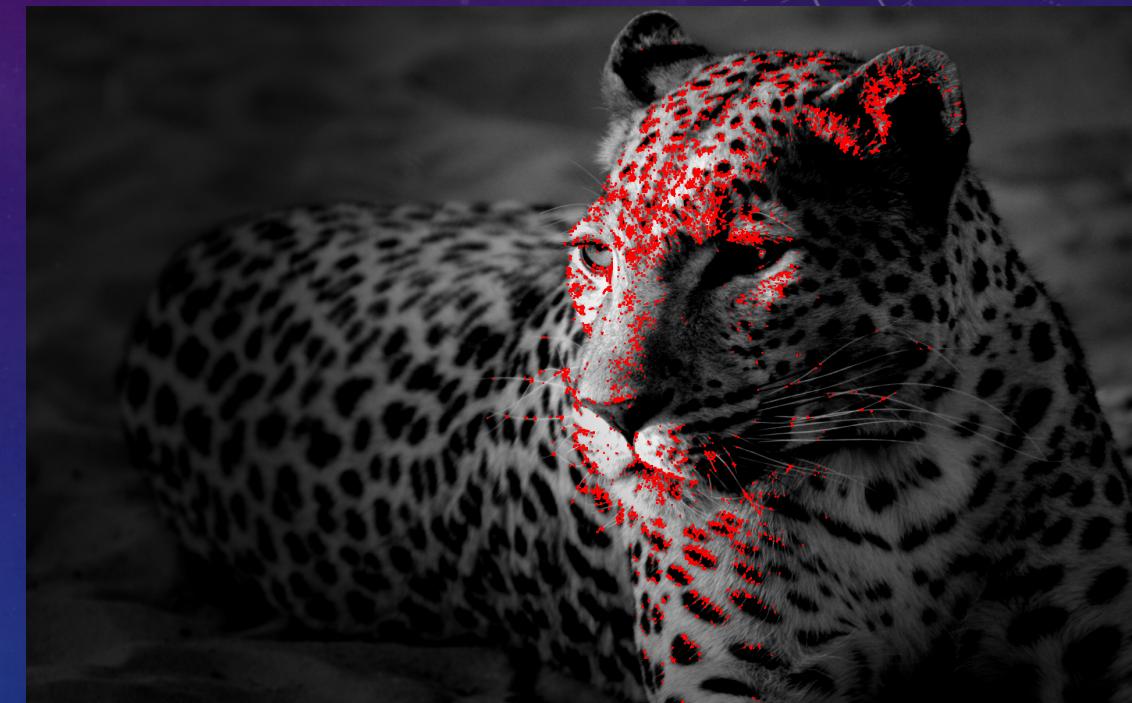
1. Compare and explain timing results
2. Look into further CUDA Optimizations

# PRELIMINARY IMAGE RESULTS



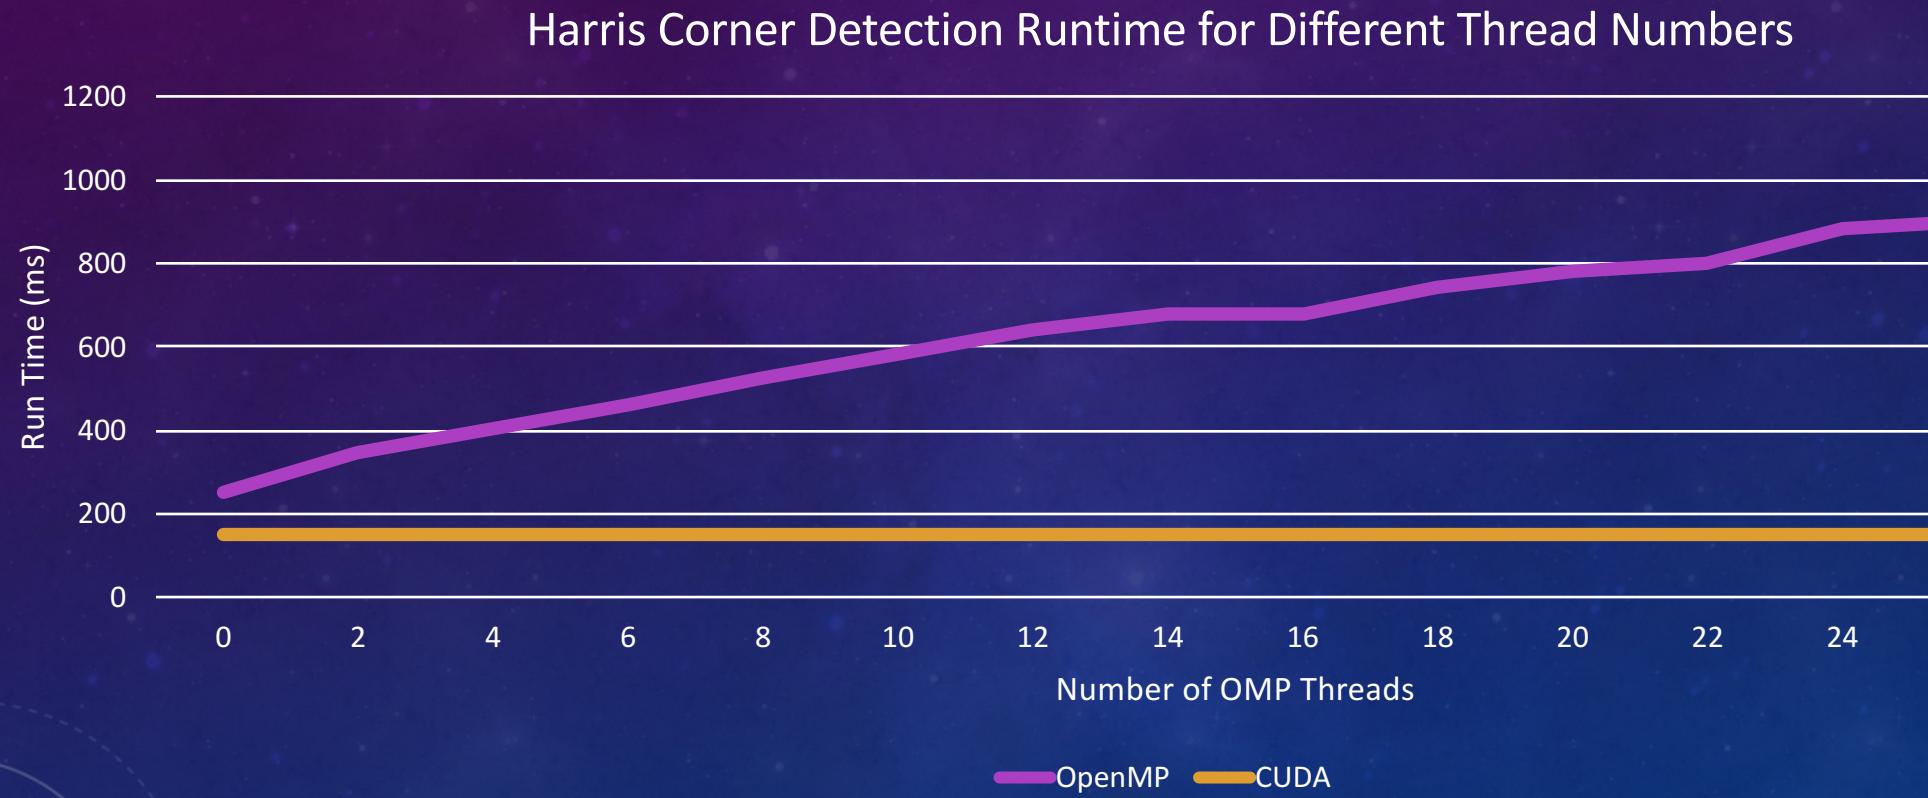
C++ Implementation

(No Non-Max Supression or Gaussian Smoothing)

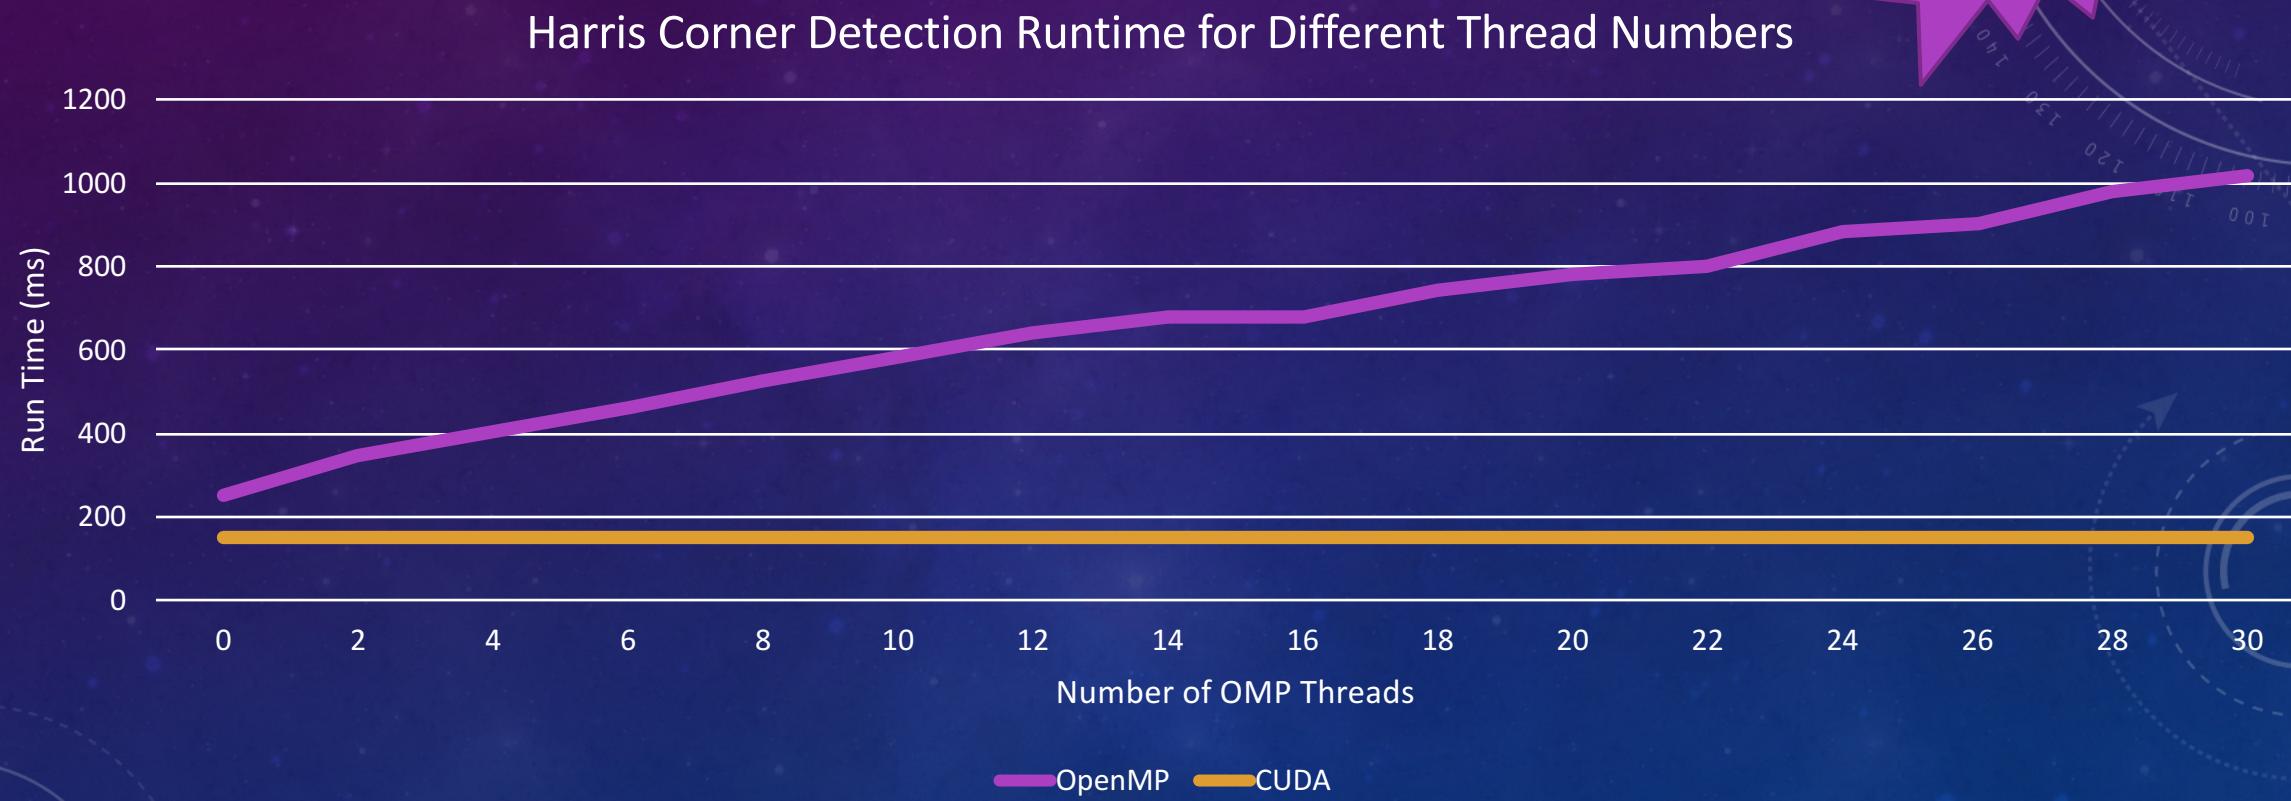


OpenCV "Ground Truth"

# PRELIMINARY TIMING RESULTS



# PRELIMINARY TIMING RESULTS



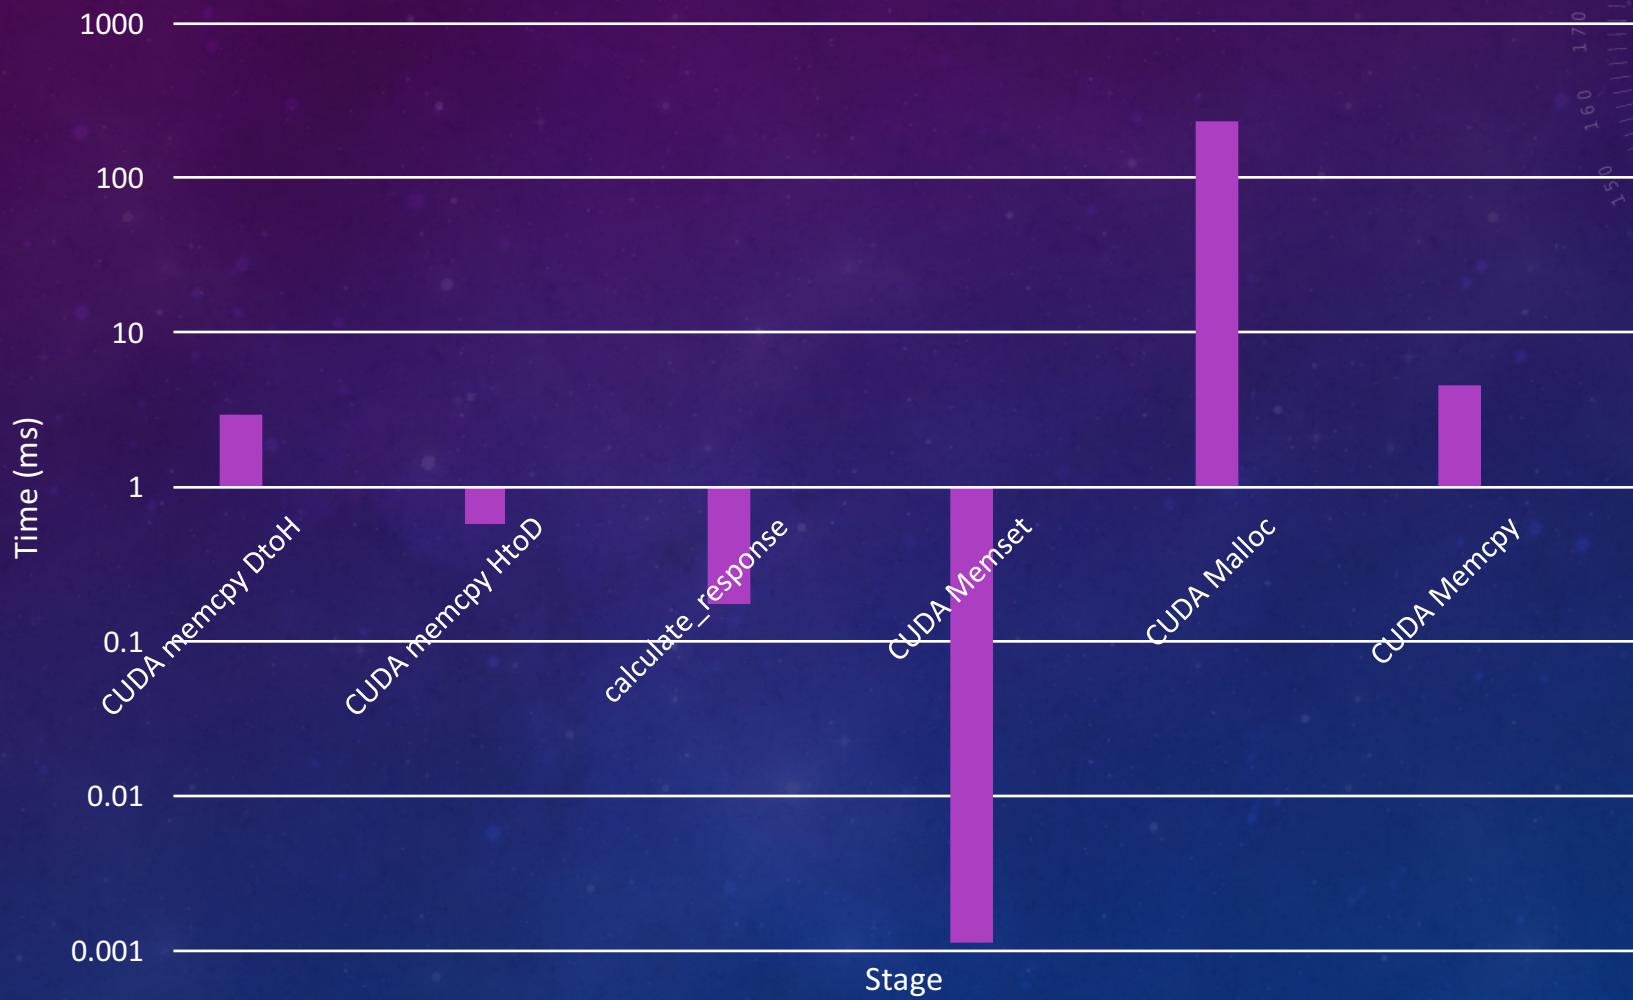
# NAÏVE CUDA IMPLEMENTATION

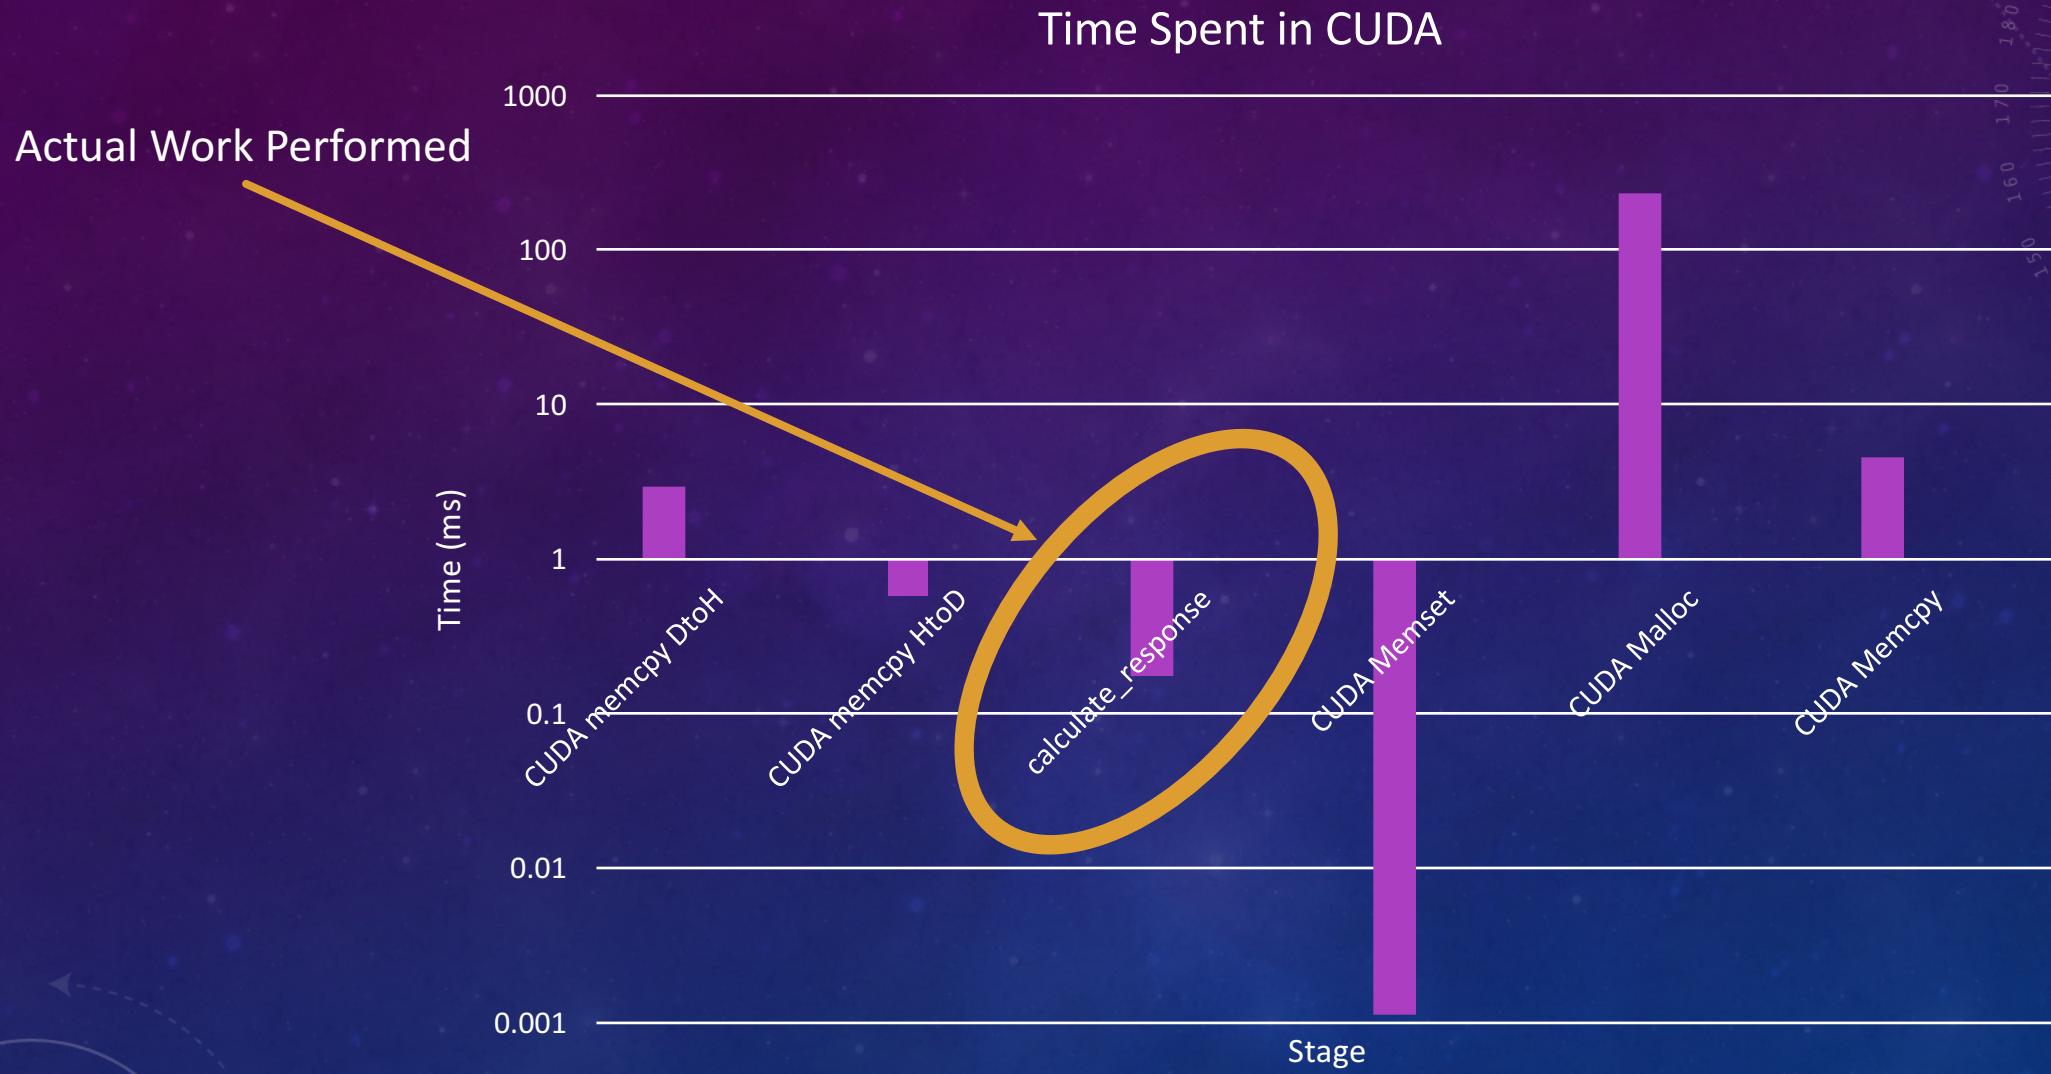
- Performs the entire pipeline except thresholding in a **single memory transfer** on the GPU
- One thread per pixel (maximize throughput)
- All data is read in and out of **GPU Global Memory**
- Thresholding is done on the CPU (serially)
- Does not perform **gaussian smoothing** or **non-max suppression**

# CUDA TIMING RESULTS

Profiling result:							
Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	79.63%	2.9945ms	1	2.9945ms	2.9945ms	2.9945ms	[CUDA memcpy DtoH]
	15.60%	586.57us	1	586.57us	586.57us	586.57us	[CUDA memcpy HtoD]
	4.74%	178.40us	1	178.40us	178.40us	178.40us	calculate_response(unsigned char*, float*, unsigned int, unsigned int)
	0.03%	1.1520us	1	1.1520us	1.1520us	1.1520us	[CUDA memset]
API calls:	96.95%	231.19ms	2	115.59ms	225.08us	230.96ms	cudaMalloc
	1.88%	14.4859ms	20	2.2429ms	928.38us	23.5575ms	cudaMemcpy
	0.54%	1.2868ms	2	643.42us	223.79us	1.0631ms	cudaFree
	0.27%	643.56us	94	6.8460us	352ns	260.15us	cuDeviceGetAttribute
	0.12%	296.21us	2	148.10us	125.57us	170.64us	cudaDeviceSynchronize
	0.09%	204.68us	1	204.68us	204.68us	204.68us	cuDeviceTotalMem
	0.08%	193.80us	1	193.80us	193.80us	193.80us	cudaLaunch
	0.03%	74.930us	1	74.930us	74.930us	74.930us	cudaMemset
	0.03%	70.885us	1	70.885us	70.885us	70.885us	cuDeviceGetName
	0.00%	9.0930us	3	3.0310us	580ns	6.8550us	cuDeviceGetCount
	0.00%	5.2110us	4	1.3020us	258ns	3.5890us	cudaSetupArgument
	0.00%	3.5600us	2	1.7800us	550ns	3.0100us	cuDeviceGet
	0.00%	2.0400us	1	2.0400us	2.0400us	2.0400us	cudaConfigureCall
	0.00%	828ns	1	828ns	828ns	828ns	cudaPeekAtLastError

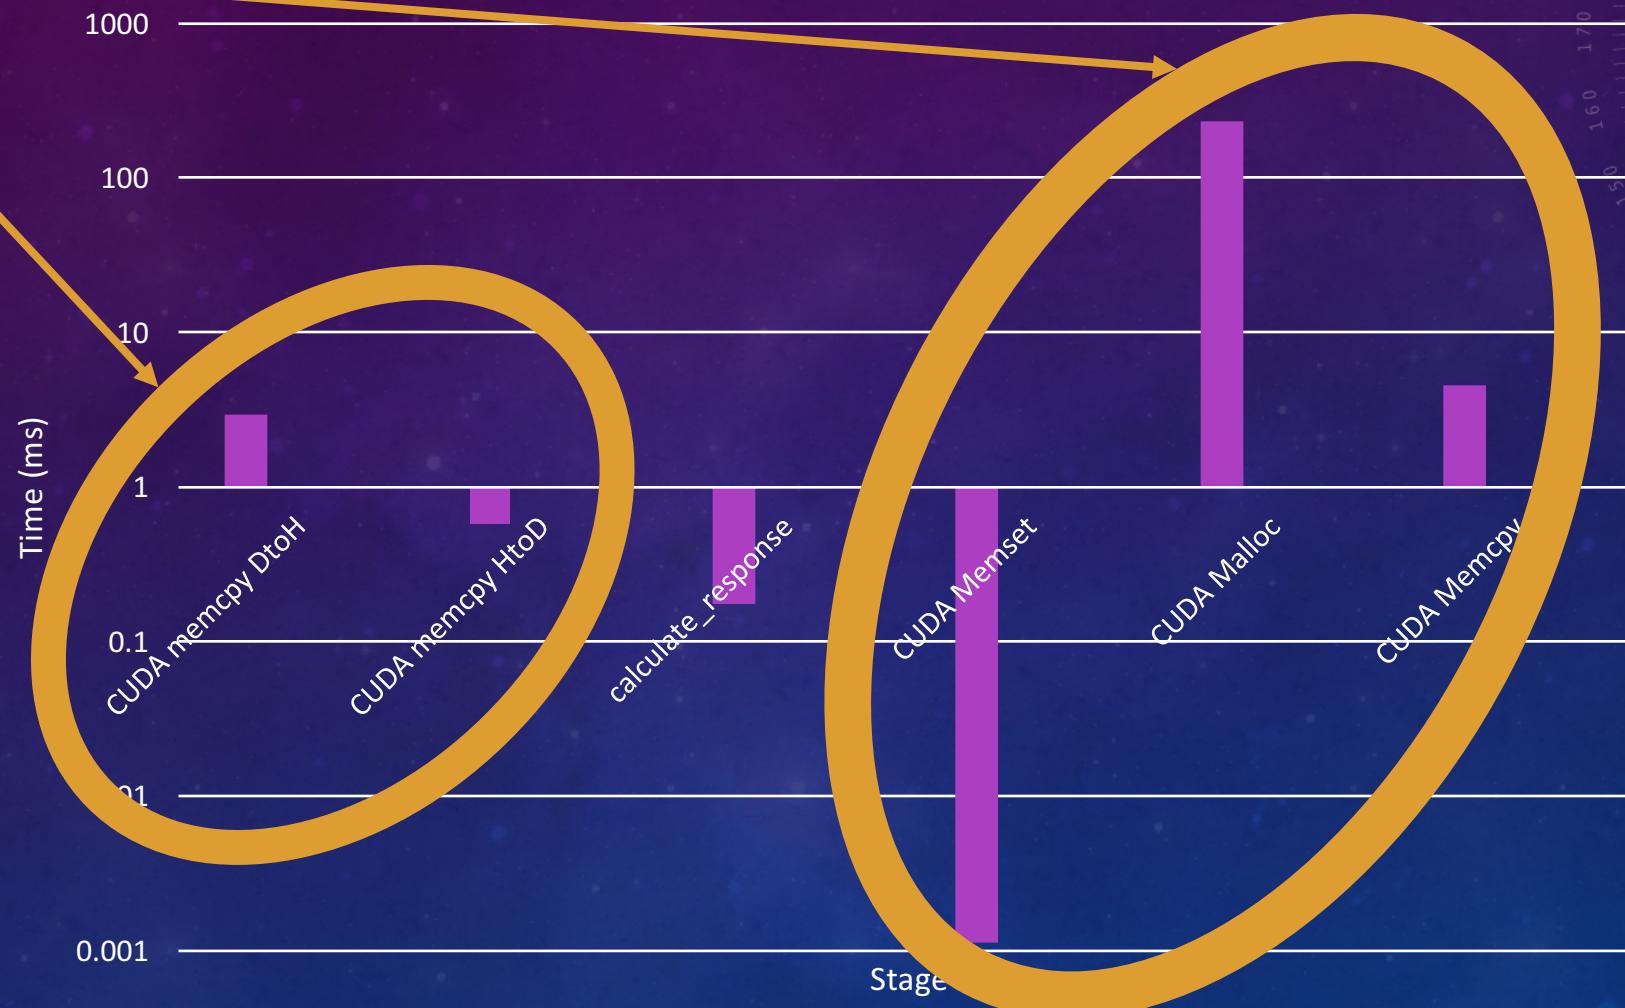
## Time Spent in CUDA





Overhead Costs!

Time Spent in CUDA



# FUTURE WORK

- Look at smarter ways of dividing up work on GPU
- Explore slow down in OpenMP
- Asynchronous transfers & convolutions
- Use of Constant + Texture Memory
- Expand Harris Corner Implementation for more robust execution