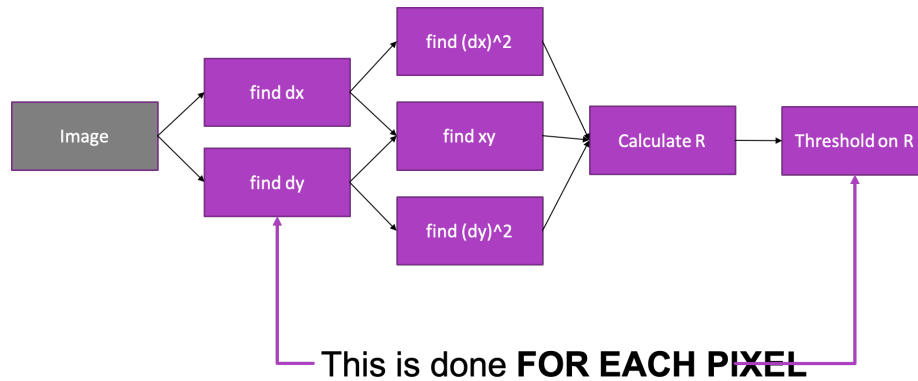


Accelerating the Harris Corner Detector through Multithreading

In this project, I implemented a basic harris corner detector in C++ and then accelerated it using OpenMP and CUDA. Image processing is an incredibly SIMD problem, lending itself incredibly well to acceleration through multithreading.




This is because all of the operations in the pipeline are done for each component pixel value, with each pixel value being performed as a stencil operation of the surrounding pixels. This means there isn't a ton of complicated control logic involved and for the most part, processing is embarrassingly parallel.

Verification

In order to verify my algorithm, I compared the output to the Harris Corner Detector implemented in OpenCV. While I don't expect the results to be exactly the same, the OpenCV method serves as a "dipstick" verification. Specifically, my implementation doesn't do the Gaussian Smoothing and non max suppression usually found in the Harris Corner Detection.



Original



Accelerated



OpenCV Ground Truth

Speedups

The first method I used for speedup was OpenMP. Most of the loop iterations in this algorithm involved iterating over pixel values. I started with a version which serialized the data from a 2D array to a 1D array to improve memory access speeds. The example below shows the calculation of gradients i_x and i_y from the input values. Each iteration of this loop is independent of one another and as a result, it can immediately be parallelized through use of OpenMP.

Calculating Gradients

```
void HarrisCorner::calculate_gradients(float* ix, float* iy, float* input, int width, int height) {
    int x, y;
    int x_minus_one, x_plus_one;
    int y_minus_one, y_plus_one;
    float dx, dy;

    // Iterate over the entire image
    #pragma omp parallel for
    for(int ii = 0; ii < width * height; ++ii) {
        x = ii % width;
        y = ii / width;

        // Break out...
```

```

        if(x <= 1 || y <= 1 || x >= width - 2 || y >= height - 2) {
            continue;
        }

        x_minus_one = y*width + (x - 1);
        x_plus_one = y*width + (x + 1);
        y_minus_one = (y - 1)*width + x;
        y_plus_one = (y + 1)*width + x;

        dx = abs(input[x_minus_one] - input[x_plus_one]);
        dy = abs(input[y_minus_one] - input[y_plus_one]);

        ix[ii] = dx;
        iy[ij] = dy;
    }
}

```

This code lends itself to CUDA. In CUDA, we spawn a single CUDA thread per pixel. From the CUDA thread/block indices and dimensions, we can calculate the x and y pixel position that each CUDA thread operates on. With this information we can also calculate the neighboring "stencil" to find the image gradients.

```

_global__ void calculate_response(unsigned char *input,
                                float *output,
                                unsigned int height,
                                unsigned int width){

    int x = blockIdx.x*TILE_SIZE+threadIdx.x;
    int y = blockIdx.y*TILE_SIZE+threadIdx.y;

    float k = .04;

    // This represents the first and last row/column of the image. sobel needs
    // a delta of at least one
    if(x <= 1 || y <= 1 || x >= width - 2 || y >= height - 2) {
        return;
    }

    // calculate the offsets for all regions of interest
    int offset = (y * width) + x;
    int x_minus_one = y*width + (x - 1);
    int x_plus_one = y*width + (x + 1);
    int y_minus_one = (y - 1)*width + x;
    int y_plus_one = (y + 1)*width + x;

    // Calculate dx, dy

```

```

float dx = input[x_minus_one] - input[x_plus_one];
float dy = input[y_minus_one] - input[y_plus_one];

// Caculate dx2, dy2, dxy
float ix2 = dx*dx;
float iy2 = dy*dy;
float ixy = dx*dy;

// Harris Corner Response Matrix:
// [ix2, ixy;
// ixy, iy2]
float itrace = ix2 + iy2;
float idet = (ix2*iy2) - (ixy*ixy);

float response = abs(idet - (k * itrace * itrace));

output[offset] = response;
}


```

One of the most important things to consider when working with the GPUs is to minimize the number of data transfer back and forth between the GPU and the host machine. The result is a function that computes the entire Harris Corner Response in a single shot per thread as opposed to pipelining by calculating larger blocks at once and combining at a higher level (i.e. computing all ix , iy values, then computing all ix^2 , iy^2 , and ixy).

Another important consideration is maximizing the "Compute per Global Memory Access" (CGMA) ratio. This is achievable by reading global memory into a local variable and running from there. However, repeated accesses are still made to global memory as a result of the repeated access incurred by the stencil operations done as part of the gradient computations. This could be further optimized by performing a SINGLE read by each thread into shared memory and then further accessing the memory from there. While some overhead is incurred by the synchronization of threads between the pre-load and execution, memory access as a whole would increase.

Unfortunately I wasn't able to implement this feature before the 8am deadline due to lack of cluster access as a result of a user using all 8 GPUs on Saturday for several hours.

Timing Results



The above graphic shows the timing results achieved by the OpenMP and CUDA implementations. OpenMP had some interesting and incredibly vexing results. On one day, iterations with different OpenMP threads showed timing increasing by 100ms for each additional pair of threads being applied. These results don't quite make sense since the additional threads should work with the embarrassingly parallel problem. At the very least it should keep results consistent, not decrease them. On a different day, the same code yielded fairly consistent timings across the board.

More interesting is the resulting timing from the CUDA implementation. I ran the CUDA code with two different tile sizes: one with a 16x16 tile and another with an 8x8 tile size. The tiles here referred to the number of threads per thread block in each dimension. With a 2D problem, I created 2D tiles.

Here we see a timing of 150ms for the 16x16 block and ~200 ms for the 8x8 block. At first glance these results don't seem too much more impressive than the original CPU implementation until you consider the breakdown of the timing.

16x16 Tiling

Activity	Type	Time(%)	Time(us)	Calls	Avg	Min	Max	Name
GPU activities:		79.60%	2.9370ms	avg all	1	2.9370ms	2.9370ms	[CUDA memcpy DtoH]
		15.54%	573.52us		1	573.52us	573.52us	[CUDA memcpy HtoD]
	Another important	4.82%	177.92us	is max	1	177.92us	177.92us	calculate_response(unsigned char*, float*, unsigned int, unsigned int)
	Access" (COM)	0.03%	1.1520us	chleevobl	1	1.1520us	1.1520us	[CUDA memset]
	API calls:	97.26%	251.71ms	However	2	125.85ms	219.04us	cudaMalloc
	Global memory	1.70%	4.4002ms	repeat	2	2.2001ms	987.36us	cudaMemcpy
	done as part of	0.49%	1.2663ms	putation	2	633.14us	217.24us	cudaFree
	performing a	510.24%	614.31us	h thread	94	6.5350us	332ns	236.43us
	accessing the	0.11%	295.32us	While	2	147.66us	120.65us	174.67us
	synchronization	0.08%	194.76us	seen the	1	194.76us	194.76us	cudaDeviceSynchronize
	in a whole world	0.06%	166.47us		1	166.47us	166.47us	166.47us
		0.03%	78.27us		1	78.27us	78.27us	78.27us
	Unfortunately	1.0.03%	68.325us	implemen	1	68.325us	68.325us	68.325us
	due to lack of	0.00%	4.8570us	a resu	4	1.2140us	310ns	3.0840us
	for several hour	0.00%	4.6280us		3	1.5420us	413ns	2.7500us
		0.00%	3.1090us		2	1.5540us	607ns	2.5020us
	Timing Result	0.00%	2.2610us		1	2.2610us	2.2610us	2.2610us
		0.00%	931ns		1	931ns	931ns	931ns


8x8 Tiling

```

==34106== Profiling application: ./release/main input/leopard.jpg out.jpg
==34106== Profiling result:
access:Type Time(%) From Time To Calls Min Avg Max Name
GPU activities: 74.41% 2.4964ms seen the 1 2.4964ms 2.4964ms 2.4964ms [CUDA memcpy DtoH]
g whole would 15.67% 532.30us 1 532.30us 532.30us 532.30us [CUDA memcpy HtoD]
9.69% 325.00us 1 325.00us 325.00us 325.00us calculate_response(unsigned char*, float*, unsigned int, unsigned int)
Unfortunately 0.03% 1.0560us implement 1 1.0560us 1.0560us 1.0560us [CUDA memset]
API calls: 96.74% 183.01ms a resu 2 91.506ms 186.09us 182.83ms cudaMalloc
for several hour 1.98% 3.7415ms 2 1.8708ms 798.96us 2.9425ms cudaMemcpy
0.64% 1.2078ms 2 603.92us 176.71us 1.0311ms cudaFree
# Timing Result 0.24% 447.58us 2 223.79us 127.17us 320.41us cudaDeviceSynchronize
0.21% 403.06us 94 4.2870us 173ns 159.67us cuDeviceGetAttribute
[[Timing]] 0.07% 140.73us 1 140.73us 140.73us 140.73us cudaLaunch
0.06% 108.94us 1 108.94us 108.94us 108.94us cuDeviceTotalMem
The above graph 0.03% 64.827us timing result 1 64.827us 64.827us 64.827us cuDeviceMemset
Implementations 0.02% 41.459us 1 41.459us 41.459us 41.459us cuDeviceGetName
On one day, I see 0.00% 5.0730us 1 5.0730us 5.0730us 5.0730us cuDeviceGetCount
100ms for each 0.00% 4.2340us of thread 4 1.0580us 235ns 2.6280us cudaSetupArgument
make sense since 0.00% 1.9170us threads 1 1.9170us 1.9170us 1.9170us cudaConfigureCall
parallel problem 0.00% 1.7280us east it 2 864ns 295ns 1.4330us cuDeviceGet
decrease them 0.00% 174.691ns by the 1 174.691ns 174.691ns 174.691ns cuDeviceGet
0.00% 174.691ns by the 1 174.691ns 174.691ns 174.691ns cuDeviceGet
0.00% 174.691ns by the 1 174.691ns 174.691ns 174.691ns cuDeviceGet

```

The function that computes the actual response (`calculate_response`) takes on the order of ~177-300 us to run, an order of magnitude less than the CPU counterpart. This shows that while the computation itself is significantly faster, a steep overhead is paid for setting up the problem. Namely, the `cudaMalloc` takes close to 97% of the time (~250ms). Looking more closely at the timing, it appears `cudaMalloc` is called twice, with the min and max values ranging WILDLY, from ~250ms as the max to a min of ~220us. At the suggestion of one of my classmates I tried "warming up" the GPU by calling a malloc outside the first timing test to see what this would do to the results. The results were drastic.



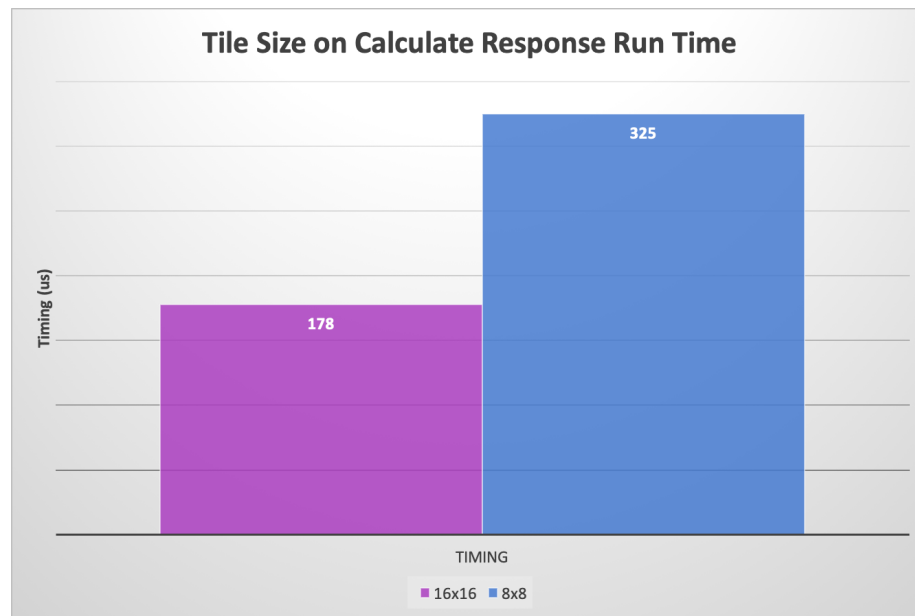
By adding in the warm up routine to prior to timing the run of the GPU, the runtime of the GPU code dropped 12ms. It's important to note here that the runtime itself did not *actually* change, just the timing of the run that did. This

is because the first call to the GPU takes a bit longer to "initialize" the GPU and bring all systems online. When we were running the benchmark, this meant that the timing data for the first malloc was causing a higher run time. This is fixed when we "warmup" the program by allocating some memory (unused later on) prior to beginning the actual computations. In essence, this experiment showed the importance of the warmup procedure when evaluating benchmarks.

This is also extended to the CPUs which are run multiple times before the first computation actually takes place.

Tiling

As part of my experiments, I explored how different sized tiles played an impact on the timing of results. The previous two nvprof screen shots show the difference in timing between different sized tiles.



The size of the tile ultimately refers to how many threads are being launched per block running on a stream multiprocessor (SM). With larger tiles, fewer blocks of more threads are launched while with smaller tiles, more blocks with fewer threads are launch. This results in differences in caching since oversubscribing the number of blocks results in more context switches between blocks which can introduce an overhead.

A tradeoff to consider (in future implementations) is the size of shared memory associated to each block for different block sizes. Since multiple blocks can run on a single SM but each SM only has a finite amount of block size, the amount

of shared memory each block can get is directly related to the number of blocks being launched.

Future Work

There were a number of components that I wanted to explore for this project that I unfortunately was not able to explore.

Shared memory:

For improving the timing on a single image, the use of shared memory would have been one of the most interesting investigations to exploit existing parallelism. While shared memory is significantly smaller than global memory, it is also significantly faster to access.

Asynchronous Access:

Asynchronous access could enable the simultaneous access and processing of data resulting in up to 2x speedups over the current implementation. This speedup comes from the ability of the GPU to begin processing data as soon as it arrives rather than waiting for all data to arrive before beginning the process it. However, this is done by taking advantage of the rasterized nature of data transfers - data is sent row by row. For certain computations which require the blending of data in different columns AND rows, temporal constraints dictate that a row must have completed transfer before processing can begin. As a result, some algorithms may need to be revised in order to accommodate this optimization.

Multiple GPUs:

Multiple GPUs can become useful in splitting up problems even further to obtain a higher degree of parallelism. My intuition says for a single image the benefits of using multiple GPUs may not outweigh the overhead in coordination and that this sort of solution is likely better suited for the larger problem sizes.

3D data (Multiple pictures)

While not an additional optimization, increasing the problem size would provide a better sense of how his algorithms scales across a number of nodes. Temporal data (such as successive frames in a video) can be incredibly compute heavy depending on the rate/size of the problem and will be able to further take advantage of the parallel hardware to solve the problems.

Code

All code is publically available on GitHub at the following link. https://github.com/drewtu2/eece5640/tree/master/final_project/harris